

Maximum Entropy Reinforcement Learning via Energy-Based Normalizing Flow

Project Presentation - Team 9

Shayan Meshkat Alsadat, Soham Karandikar,
Aman Chandak, and Abhinav V Pillai

Arizona State University

November 23, 2025

Problem Statement

- Existing MaxEnt RL methods for continuous action spaces rely on actor-critic frameworks, leading to complex dual optimization and potential misalignment between actor and critic networks.
- Sampling actions directly from the underlying energy-based model of the soft Q-function is computationally expensive, often necessitating a separate actor network for efficient sampling.
- Calculating the soft value function involves an intractable integral, forcing reliance on Monte Carlo approximations that introduce estimation errors and variance.

- **Maximum Entropy RL:** Integrates policy entropy as intrinsic reward
 - Improves exploration-exploitation balance
 - Enhanced robustness in continuous control
- **Key Challenges:**
 - Learning soft Q-function requires costly MCMC sampling
 - Soft value calculation needs computationally infeasible integration
 - Actor-critic separation introduces optimization errors
- **Authors' Solution:** MEow - A unified framework using Energy-Based Normalizing Flows

Background: Maximum Entropy RL

MaxEnt Objective

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\rho_{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t (r(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot|s_t))) \right]$$

- Optimal policy expressed via soft Q-function and soft value function:

$$\pi^*(a|s) = \exp \left(\frac{1}{\alpha} (Q^*(s, a) - V^*(s)) \right)$$

$$V^*(s) = \alpha \log \int_{\mathcal{A}} \exp \left(\frac{1}{\alpha} Q^*(s, a) \right) da$$

- Problem:** Computing $V^*(s)$ requires intractable integration

Two Estimation Methods:

- SQL:** Importance sampling

$$V(s) \approx \alpha \log \frac{1}{N} \sum_{i=1}^N \exp \left(\frac{1}{\alpha} Q(s, a_i) \right)$$

- SAC:** Policy approximation

$$V(s) \approx \mathbb{E}_a [Q(s, a) - \alpha \log p(a|s)]$$

Can we do unified approach with exact computation?!

Energy-Based Normalizing Flows

- **Normalizing Flows:** Invertible transformations for density modeling

$$p(x; \theta) = p_Z(f_\theta^{-1}(x)) \prod_{i=1}^L \left| \det \frac{\partial f_{\theta,i}}{\partial f_{i-1}} \right|^{-1}$$

- **Energy-Based View:** Decompose Jacobian by linear/nonlinear layers

$$p(x; \theta) = \frac{1}{Z(\theta)} \exp(-E(x; \theta))$$

where:

- Energy: $E(x; \theta) = -\log p_Z(f_\theta^{-1}(x)) - \sum_{i \in \mathcal{I}_N} \log \left| \det \frac{\partial f_i}{\partial f_{i-1}} \right|$
- Normalizing constant: $Z(\theta) = \prod_{i \in \mathcal{I}_L} \left| \det \frac{\partial f_i}{\partial f_{i-1}} \right|$
- **Key Property:** $Z(\theta)$ is independent of input x !

M**E**ow Framework: Core Idea

Is there a way to calculate optimal value function without approximating it with Monte-Carlo methods?

$$V^*(s) = \alpha \log \int_{\mathcal{A}} \exp \left(\frac{1}{\alpha} Q^*(s, a) \right) da$$

It turns out we could!!

Unified Policy Representation

$$\pi(a|s; \theta) = \frac{1}{Z(s; \theta)} \exp \left(\frac{1}{\alpha} Q(s, a; \theta) \right)$$

Theorem (Exact Soft Value)

$$V(s; \theta) = \alpha \log \int_{\mathcal{A}} \exp \left(\frac{1}{\alpha} Q(s, a; \theta) \right) da = -\alpha \log Z(s; \theta)$$

- Single training objective (no actor-critic split)
- Exact value computation (no approximation error)
- Efficient sampling via normalizing flow

Contributions: MEow Framework

① Unified Policy Representation

- Single model for both Q-function and policy
- Eliminates actor-critic separation

② Exact Soft Value Computation

- No Monte Carlo approximation needed
- $V(s) = -\alpha \log Z(s; \theta)$ computed exactly

③ Enhanced Training Techniques

- Learnable Reward Shifting (LRS) for numerical stability
- Shifting-Based Clipped Double Q-Learning (SCDQ) for overestimation bias

④ Superior Empirical Performance

- MuJoCo benchmarks
- Omniverse Isaac Gym robotic tasks

Training Objective:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s')} \left[(Q(s, a; \theta) - r - \gamma V(s'; \theta'))^2 \right]$$

- Single loss function
- Exact $V(s')$ via $-\alpha \log Z(s'; \theta')$
- Target network with Polyak averaging

Inference:

- **Stochastic:** Sample $z \sim p_Z(z)$, compute $a = f_\theta(z|s)$
- **Deterministic:**

$$a^*(s) = f_\theta(\mathbb{E}[z]|s)$$

Multi-modal Support:

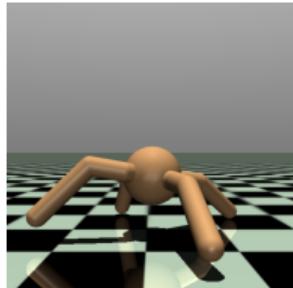
- Normalizing flows can represent complex, multi-modal action distributions

Training Algorithm

```
1 Init:  $\theta, \theta' \leftarrow \theta, \alpha, \gamma, \tau, \beta, g_\theta, b_\theta^{(1)}, b_\theta^{(2)}$ , buffer  $\mathcal{D}$     // Initialize networks, targets, replay buffer
2 for each step do
3     Sample  $z \sim p_z$                                               // Sample latent noise
4      $a_t \leftarrow g_\theta^{-1}(z | s_t)$                                 // Invert flow to get action
5      $s_{t+1} \sim p^T(\cdot | s_t, a_t)$                                 // Environment transition
6     Store  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$                           // Add to replay
7     Sample batch  $(s_t, a_t, r_t, s_{t+1}) \sim \mathcal{D}$                 // Replay sampling
8      $Q_\theta = \alpha \log \left[ p_z(g_\theta(a_t | s_t)) \prod_{i \in \mathcal{S}_\ell} |\det(J_{g_\theta^i})| \right]$  // Flow-based Q
9      $V_{\theta'} = -\alpha \log \prod_{i \in \mathcal{S}_\ell} |\det(J_{g_{\theta'}^i})|$           // Flow-based V
10     $Q_\theta^{(k)} = Q_\theta + b_\theta^{(k)}$                                 // Twin critics
11     $V_{\theta'}^{clip} = V_{\theta'} + \min(b_{\theta'}^{(1)}, b_{\theta'}^{(2)})$         // Clipped target
12     $y_t = r_t + \gamma V_{\theta'}^{clip}(s_{t+1})$                             // TD target
13     $\mathcal{L} = \frac{1}{2}(Q_\theta^{(1)} - y_t)^2 + \frac{1}{2}(Q_\theta^{(2)} - y_t)^2$  // Bellman loss
14     $\theta \leftarrow \theta + \beta \nabla_\theta \mathcal{L}$                            // Gradient update
15     $\theta' \leftarrow (1 - \tau)\theta' + \tau\theta$                             // Soft update
```

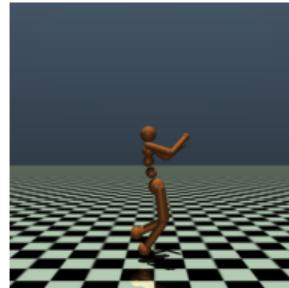
Results: MuJoCo Benchmarks Overview

Ant-v4



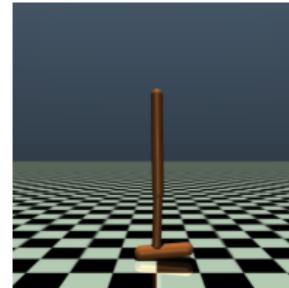
State dim: 27 | Action dim: 8

Humanoid-v4



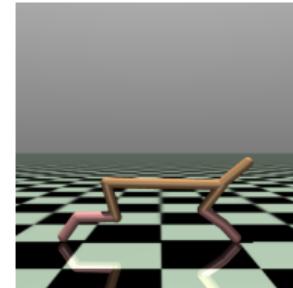
State dim: 376 | Action dim: 17

Hopper-v4



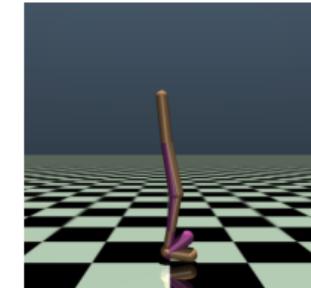
State dim: 11 | Action dim: 3

HalfCheetah-v4



State dim: 17 | Action dim: 6

Walker2d-v4



State dim: 17 | Action dim: 6

Tasks: Ant, Humanoid, Hopper, HalfCheetah, Walker2d

Setup: 1.5M - 5M timesteps depending on task, 5 seeds per method

State/Action Dimensions:

- Hopper: 11/3 (simplest)
- Humanoid: 376/17 (most challenging)

MuJoCo Results: Simpler Tasks

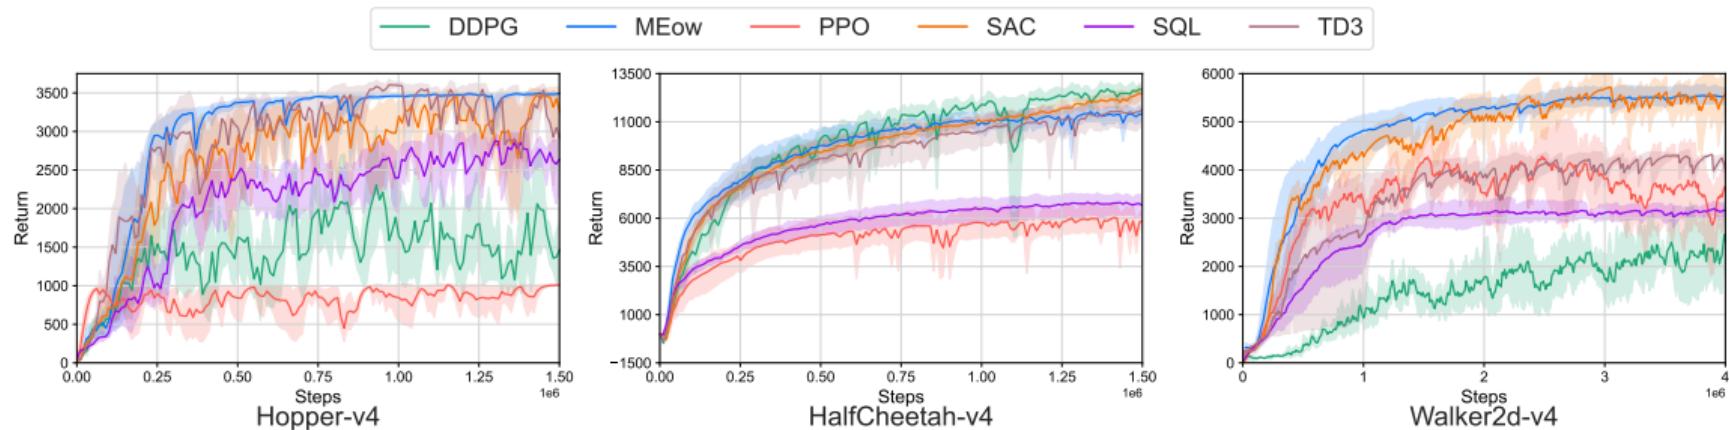


Figure: Results for Hopper-v4, HalfCheetah-v4 and Walker2d-v4.

Key Observations:

- MEow achieves comparable performance to SAC across all three tasks
- Competitive learning speed and stable convergence
- Outperforms DDPG, TD3, and PPO consistently

MuJoCo Results: Simpler Tasks

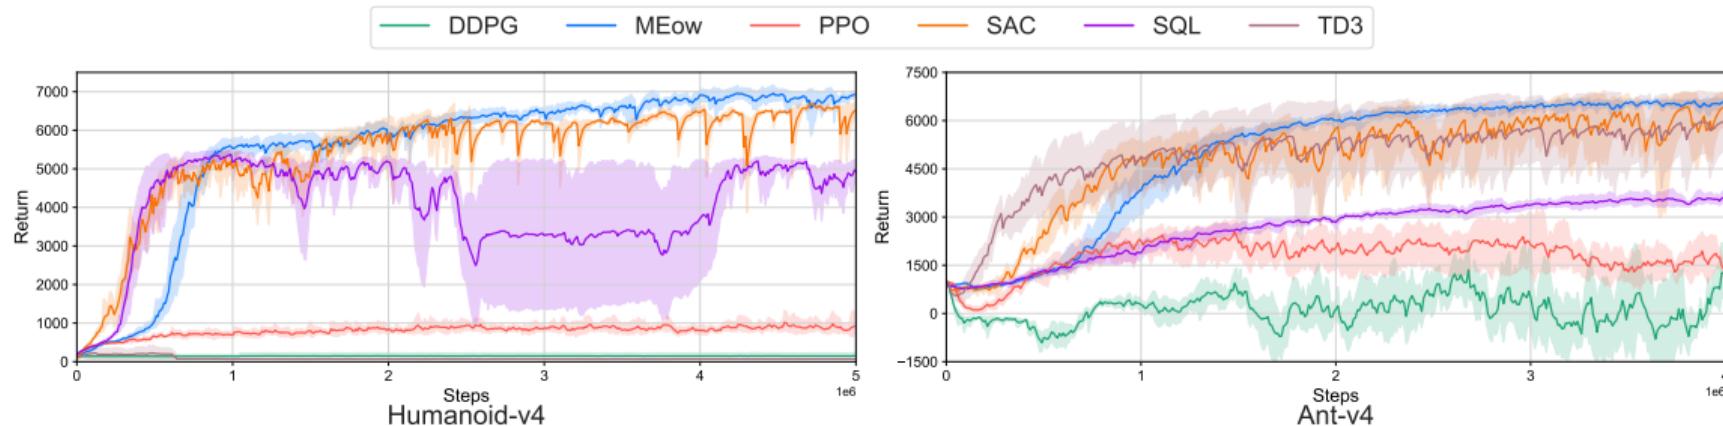


Figure: Results for Humanoid-v4 and Ant-v4.

Key Observations:

- Notably smoother curves with fewer evaluation spikes
- Reduced variance suggests more stable optimization
- Superior sample efficiency and training stability across both high-dimensional environments

Results: Omniverse Isaac Gym - Overview

6 Robotic Tasks: Designed for real-world deployment scenarios

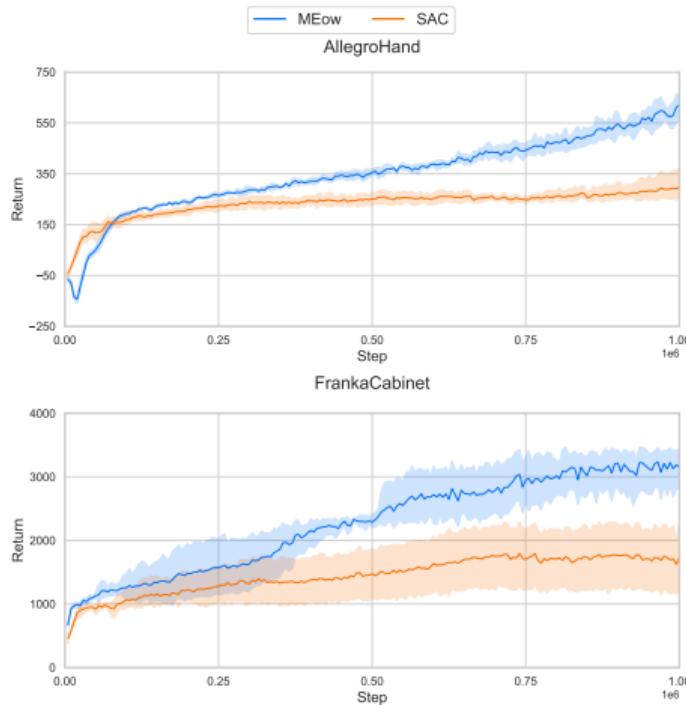
Key Results Summary

- MEow **consistently outperforms SAC** across all tasks
- Substantial improvements in sample efficiency
- Better asymptotic performance
- Smoother learning curves

Isaac Gym: Manipulation Tasks

AllegroHand - Dexterous Manipulation:

- MEow: **550** return vs SAC: 350
- **28% improvement**
- 16 actuators, 88-dimensional state space
- Better handling of **multi-modal action distributions**



FrankaCabinet - Robotic Arm Manipulation:

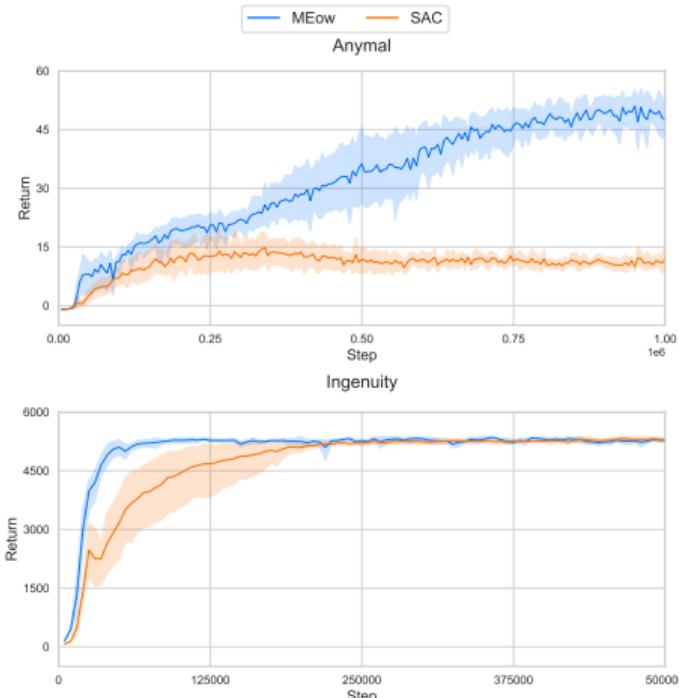
- MEow: **3000** return vs SAC: 1500
- **29% improvement**
- Contact-rich interactions with cabinet doors
- Exact value estimation particularly beneficial

Figure: Learning curves for manipulation tasks

Isaac Gym: Locomotion Tasks

ANYmal - Quadruped Locomotion:

- MEow: **45** return vs SAC: 15
- **17% improvement**
- Notably smoother learning curves
- Stable locomotion emerges **faster**
- Reduced optimization variance



Ingenuity - Helicopter Flight Control:

- MEow: **5000** return vs SAC: 5000
- comparable performance
- Challenging underactuated system
- Performance gap emerges **early in training**
- Better exploration efficiency

Figure: Learning curves for locomotion tasks

Summary: Bottlenecks and Conclusions

MEow Bottlenecks

- 2× slower than SAC during training
- Slow real-time interaction for high-frequency control
- Sensitive to hyperparameter tuning

MEow Key Contributions

- Unified objective (no actor–critic split)
- Exact soft value (no MC approximation)
- Improved stability (LRS + SCDQ)
- Strong empirical results

Application Challenges:

- Poor generalization to unseen states
- Limited hardware/deployment readiness
- Difficulty scaling to high-dimensional actions

Advantages:

- No actor–critic inconsistency
- Supports multi-modal policies
- Efficient, deterministic sampling

Future Directions:

- Trade-off between expressiveness and speed

Thank You!

Questions?

Original paper: [Maximum Entropy Reinforcement Learning via Energy-Based Normalizing Flow](#),
NeurIPS 2024

Contact: {smeshka1, skarand4, achan160, avpilla2}@asu.edu